


Lec Mon

Monday, October 15, 2018 3:58 PM

K Nearest Neighbors

Motivation: How would we color the blank circles?



Question: How are the boundaries different in KNN and decision tree?

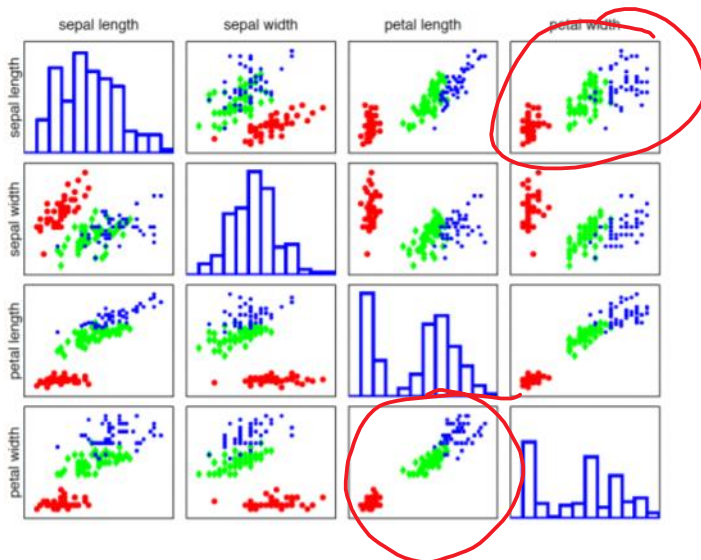
Motivation: (Computer Vision) How can we classify what type of flower it is from a photo?

- KNN was widely used in CV before deep learning was around

Example data:

Fisher's Iris Data				
Sepal length	Sepal width	Petal length	Petal width	Species
5.1	3.5	1.4	0.2	<i>I. setosa</i>
4.9	3.0	1.4	0.2	<i>I. setosa</i>
4.7	3.2	1.3	0.2	<i>I. setosa</i>
4.6	3.1	1.5	0.2	<i>I. setosa</i>
5.0	3.6	1.4	0.2	<i>I. setosa</i>
5.4	3.9	1.7	0.4	<i>I. setosa</i>
4.6	3.4	1.4	0.3	<i>I. setosa</i>
5.0	3.4	1.5	0.2	<i>I. setosa</i>
4.4	2.9	1.4	0.2	<i>I. setosa</i>
4.9	3.1	1.5	0.1	<i>I. setosa</i>

Which features make a good separator?



Since the clusters are physically more separated

Now onto the algorithm:

Nearest Neighbors: The basic version

- training examples are vectors x_i associated with label y_i
- Learning: just store training examples
- Prediction: for a new example x ,
 - find the training example x_i that is closest to x
 - predict the label of x to be the label y_i associated with x_i

K-Nearest Neighbors

Same as above, but:

- find the **k closest** training examples to x
- construct the label of x using these k points

Issues in designing KNN modeling

- how to choose k ?
- how to define distance?
- how to aggregate the information from the k neighbors and make the prediction?

What kind of problems can KNN solve?

- classification: every neighbor votes on the label, predict the most frequent label among the neighbors
- regression: predict the mean value

Distance - how to measure it?

- **Euclidean distance**

$$\|x_1 - x_2\|_2 = \sqrt{\sum_{i=1}^n (x_{1,i} - x_{2,i})^2}$$

- **Manhattan distance**

$$\|x_1 - x_2\|_1 = \sum_{i=1}^n |x_{1,i} - x_{2,i}|$$



- **L_p -norm**

- Euclidean = L_2
- Manhattan = L_1

$$\|x_1 - x_2\|_p = \left(\sum_{i=1}^n |x_{1,i} - x_{2,i}|^p \right)^{\frac{1}{p}}$$

$$p \in \mathbb{Z}^+$$

- Most common distance is the **Hamming Distance**

- number of bits that are different
- or number of features that have different value
- Ex:
 - $x_1: \{\text{shape}=\text{triangle}, \text{color}=\text{red}, \text{location}=\text{left}, \text{orientation}=\text{up}\}$
 - $x_2: \{\text{shape}=\text{triangle}, \text{color}=\text{blue}, \text{location}=\text{left}, \text{orientation}=\text{down}\}$
 - Hamming Distance = 2
- we could apply weights to features if some are more important

How to find K?

Recap: Tuning parameters using a validation set

- Split data into:
 - Training Data
 - to train the model
 - Test Data
 - to evaluate the model
 - Development Data
 - to tune hyperparameters
- try a number of values for k. For each value,
 - train a model using training set
 - evaluate performance on development set
 - choose the value with best performance on dev set


Tips/good practice

- In general, use an odd k to avoid ties
- **Feature Normalization** to prevent larger-valued features from dominating others in distance calculations
 - normalize data to have zero mean and unit variance in each dimension
 - how to deal with categorical features? usually we turn categorical into numerical features

decision boundary for KNN

- is KNN explicitly building a function? no

It can be represented as a **Voronoi Diagram**:



Instance based learning

- a class of learning methods involving
 - learning: storing examples
 - classifying based on existing, similar examples
- most computation is performed only at prediction time
 - like open vs closed book exams

Advantages

- training time ~ 0
- easy to update the "model", since we just add the new data points
- can learn complex decisions

Disadvantages

- hard to represent outliers
- stores all examples, model can be very big
- predicting new examples takes lots of time
- curse of dimensionality - things can go wrong in high-dimensional spaces
 - what's intuitive for 2-3 dimensions may not be for more dimensions
- because of the dimensionality curse, "neighborhood" becomes very large - distances become large, and certain features will dominate over less relevant ones